

Mango: A Python Library for Parallel Hyperparameter Tuning

Sandeep Singh Sandha¹, Mohit Aggarwal², Igor Fedorov², Mani Srivastava¹


¹ University of California, Los Angeles

² Arm Research

<https://github.com/ARM-software/mango>





Tuning Hyperparameters



Approaches: Random search, Grid search, Bayesian optimization, Gradient-based, Evolutionary and Population.

ML algorithms are often **very sensitive** to choice of hyperparameters.

Example: SVM (C and gamma parameters)



Tuning Hyperparameters

Start

Tuning ..

Re-Tuning ..

Re-Tuning ..


Approaches: Random search,
Grid search, Bayesian
optimization, Gradient-based,
Evolutionary and Population.

Hyperparameter tuning is a tedious task and may involve a significant amount of software engineering.


ML algorithms are often very sensitive to choice of hyperparameters

Example: SVM (C and gamma parameters)

Grid Search



Random Search

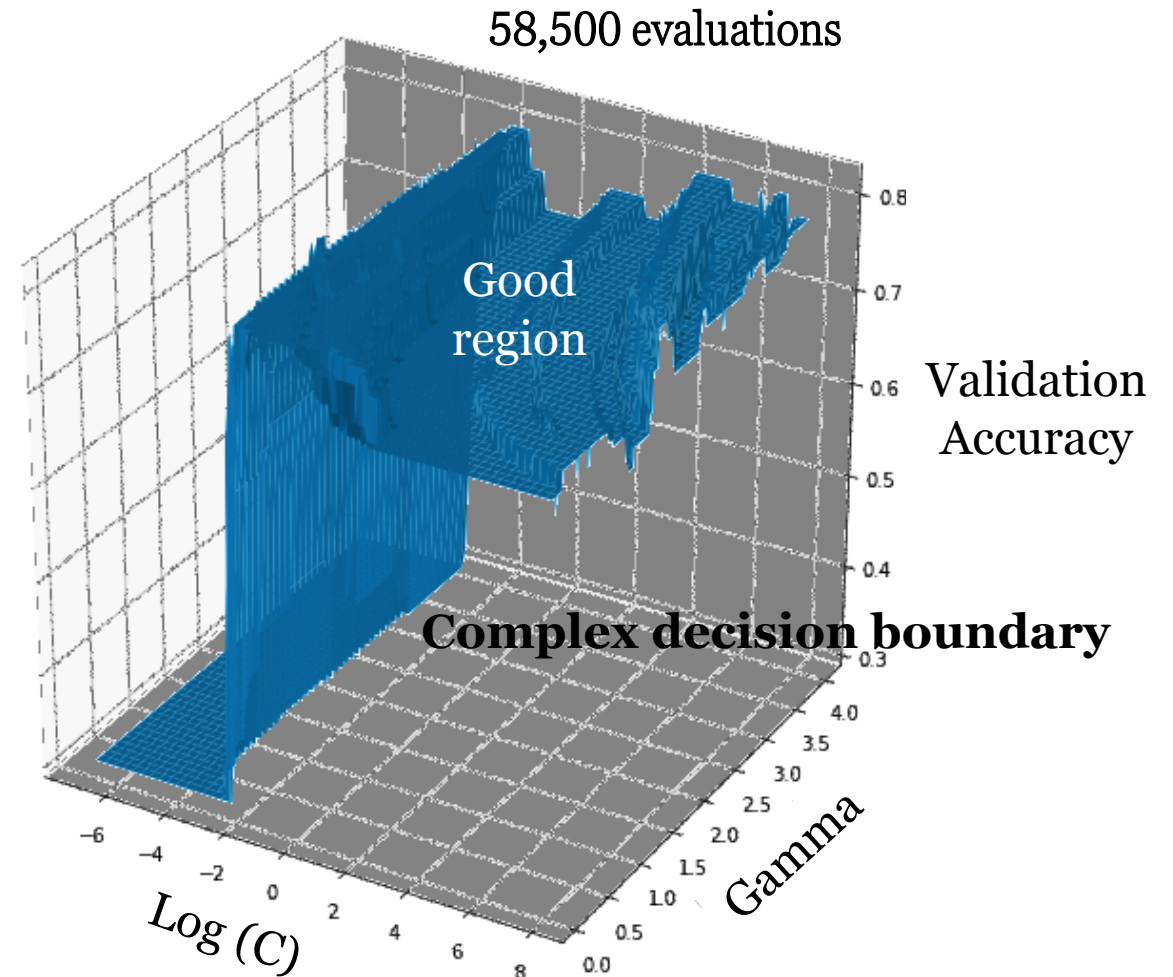


Hyperparameters: An Example

- SVM classifier for Iris Dataset
 - 150 examples, used for training
 - 50 examples used for validation
 - Used 2 features out of 4
- Hyperparameters:
 - kernel: rbf
 - gamma: $[0.1, 4]$
 - C: $[10^{-6}, 10^8]$
- Best accuracy: ?

Hyperparameters: An Example

- SVM classifier for Iris Dataset
 - 150 examples, used for training
 - 50 examples used for validation
 - Used 2 features out of 4
- Hyperparameters:
 - kernel: rbf
 - gamma: $[0.1, 4]$
 - C: $[10^{-6}, 10^8]$
- Best accuracy: 82%



Hyperparameters: An Example

XGBoost classifier parameter search space: ~1 million with coarse grids.

```
param_space = {  
    'n_estimators': range(10, 2001, 100),  
    'max_depth': range(1, 15),  
    'reg_alpha': loguniform(-3, 6),  
    'booster': ['gbtree', 'gblinear'],  
    'colsample_bylevel': uniform(0.05, 0.95),  
    'colsample_bytree': uniform(0.05, 0.95),  
    'learning_rate': loguniform(-3, 3),  
    'reg_lambda': loguniform(-3, 6),  
    'min_child_weight': loguniform(0, 2),  
    'subsample': uniform(0.1, 0.89)  
}
```

Hyperparameters: An Example

XGBoost classifier parameter search space: ~1 million with coarse grids.

```
param_space = {  
    'n_estimators': range(10, 2001, 100),
```

Grid Search may not be optimal and often is very time consuming.
Does random search work?

```
    'colsample_bytree': uniform(0.05, 0.95),  
    'learning_rate': loguniform(-3, 3),  
    'reg_lambda': loguniform(-3, 6),  
    'min_child_weight': loguniform(0, 2),  
    'subsample': uniform(0.1, 0.89)  
}
```

What is Mango?

- Mango
 - Intelligent parallel search algorithms
 - Modular design: Ease of usability
 - Abstracts are compatible with Sklearn
- We Focus on classical ML algorithms
 - Designed for a production cluster
- Why Mango?
 - Existing widely used libraries: Hyperopt, Auto-sklearn, Auto-WEKA
 - Existing libraries were not designed for production cluster
 - Many have **missing feature**: Fault tolerance, Job scheduling challenges, Parallel search, Compatibility.
 - Assumes specific kind of platform/workers are running or are totally manual.

[1] James Bergstra, Dan Yamins, and David D Cox, “Hyperopt: A python library for optimizing the hyperparameters of machine learning algorithms,” in Proceedings of the 12th Python in science conference. Citeseer, 2013

[2] Chris Thornton, Frank Hutter, Holger H Hoos, and Kevin Leyton-Brown, “Auto-weka: Combined selection and hyperparameter optimization of classification algorithms,” in Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining, 2013

[3]. Auto-Sklearn, “Manual parallel spawning,” <https://automl.github.io/auto-sklearn/master/manual.html>

What is Mango?

- Implementation of state-of-the-art parallel search optimizers
- Available opensource
- Continually tested on production scale datasets and is improved

<https://github.com/ARM-software/mango>

Mango: Simple Example

```
1  from mango import scheduler, Tuner
2
3  param_space = dict(x=range(-10,10))
4
5  @scheduler.serial
6  def objective(x):
7      return x * x
8
9  tuner = Tuner(param_space, objective)
10 results = tuner.minimize()
11
12 print(results["best_params"],
13       results["best_objective"])
```

Search Space: 3

Objective Function: 5-7

Run Mango Tuner: 9-10

Results: 12-13

Mango Abstractions: Optimizer

- Multi-armed bandit bayesian optimizer
- Surrogate function: Gaussian Process (**Internal approximator**)
- Acquisition function: Created from Surrogate function (**Good regions to explore**)

Mango Abstractions: Optimizer

- Multi-armed bandit bayesian optimizer
- Surrogate function: Gaussian Process (**Internal approximator**)
- Acquisition function: Created from Surrogate function (**Good regions to explore**)

Simplified form of optimizer:

1. **Sample the parameter search** space using Monte Carlo.
2. **Find a batch** of next promising hyperparameter to evaluate.
 1. Maximize the acquisition function (exploration vs. exploitation).
 2. Implemented two batch selection approaches: **hallucination, clustering**.
3. **Update the acquisition function based on results.**

Mango Abstractions: Optimizer

- Multi-armed bandit bayesian optimizer
- Surrogate function Gaussian Process
- Acquisition function

Evaluate promising regions in parallel

Simplified form:

Adaptive sampling

1. Sample the parameter search space using Monte Carlo.

2. Find a batch of promising regions to evaluate.

1. Maximize the expected value of the acquisition function (vs. exploitation)

2. Implemented two batch selection approaches: hallucination, clustering.

3. Update the acquisition function

Adaptive exploration

Fault tolerance

Mango Abstractions: Optimizer


- Multi-armed bandit bayesian optimizer
- Surrogate function: Gaussian Process
- Acquisition function

Simplified for

1. Sample
2. Find a bound
3. Update

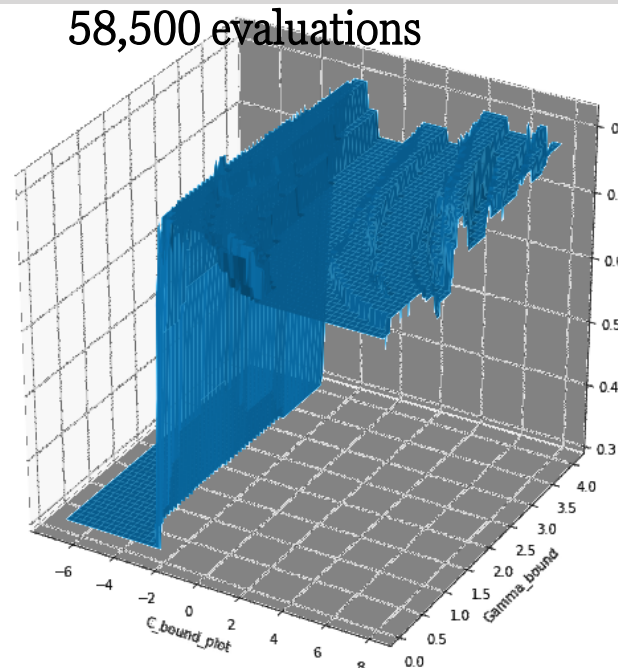
1. Maximize
2. Improve

50 evaluations



Mango Approximator

58,500 evaluations




Dense Grid Search

e.
(exploitation)
ation, clustering.

Mango Abstractions: Optimizer

Visualizing clustering approach



Mango Abstractions: Search Space

- Example: XGBoost Classifier

```
param_space = {  
    'n_estimators': range(10, 2001, 100),  
    'max_depth': range(1, 15),  
    'reg_alpha': loguniform(-3, 6),  
    'booster': ['gbtree', 'gblinear'],  
    'colsample_bylevel': uniform(0.05, 0.95),  
    'colsample_bytree': uniform(0.05, 0.95),  
    'learning_rate': loguniform(-3, 3),  
    'reg_lambda': loguniform(-3, 6),  
    'min_child_weight': loguniform(0, 2),  
    'subsample': uniform(0.1, 0.89)  
}
```

- Search space

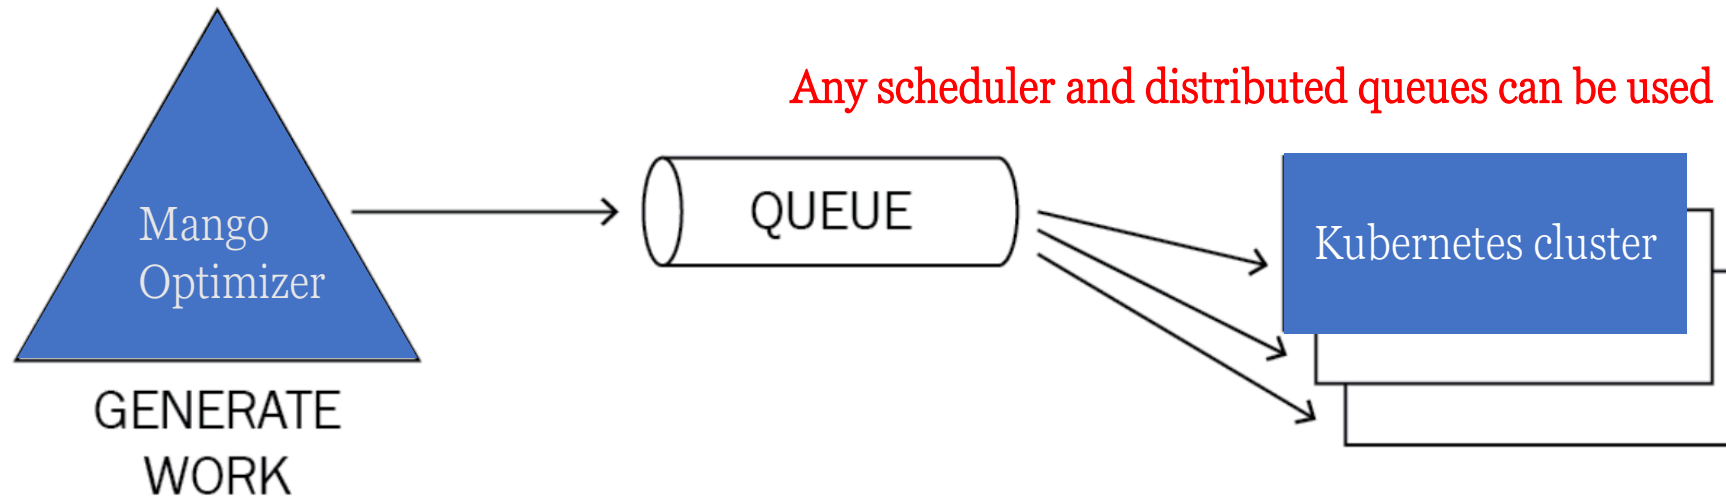
- We support all 70+ distributions of `scipy.stats`
- Python constructs lists, dicts, range.
- **Compatible with sklearn** grid search
- **Continuous**: distributions
- **Discrete**: lists, strings
- New distributions can be easily defined.

- Example: SVM Classifier

```
param_space = {  
    'kernel': ['rbf', 'sigmoid'],  
    'gamma': uniform(0.1, 4),  
    'C': loguniform(-7, 8)  
}
```

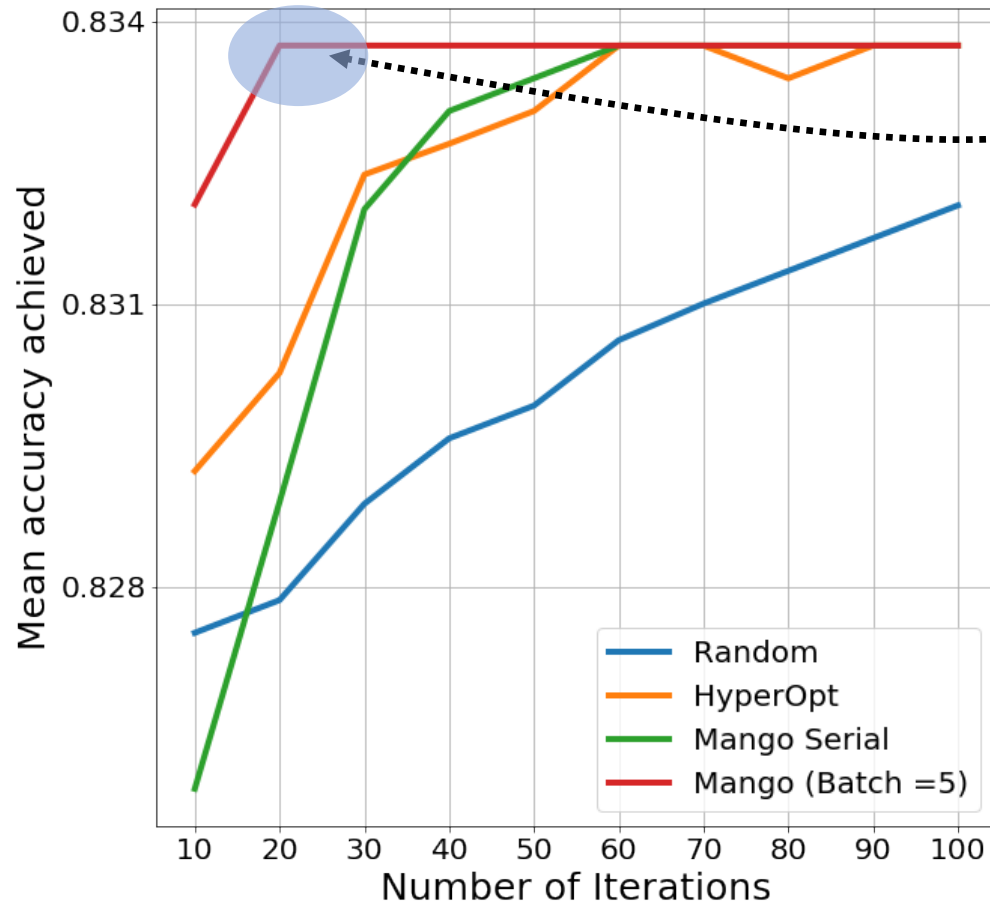

Mango Abstractions

- Objective Function & Scheduler
 - Objective function defined by user
 - Scheduler abstractions are kept separate from optimizer
- Objective Function
 - Objective function receives an hyperparameter to evaluate.
 - In case of failure, it can return empty results



Celery workflow used in production

Mango Experiments




Batch Mango is much faster


Settings: **SVM Classifier**

- Iris dataset
- Mean of 20 experiments
- HyperOpt: Another widely used library
- Mango in batch setting (hallucination)

Mango Approx.



Mango Experiments



Batch Mango is much faster


Settings: **XGBoost** Classifier

- Wine dataset
- Mean of 20 experiments

[1] Chen, Tianqi, and Carlos Guestrin. "Xgboost: A scalable tree boosting system." Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining. 2016.

[2] <https://github.com/ARM-software/mango/tree/master/examples/>

Mango Experiments




Batch Mango is much faster

Settings: **Branin function**

- Mixed discrete and continuous variables
- Mean of 20 experiments

Mango Experiments



Mango competes with the existing state-of-the-art libraries.

- Mixed discrete and continuous variables
- Mean of 20 experiments

Conclusion

- Parallel search speeds up the task of hyperparameter tuning.
- Mango achieves comparable performance to existing approaches:
 - Provide rich abstractions for search space.
 - Intelligent parallel search algorithms.
 - Flexibility to the choice of scheduler.
 - Designed with the ease of usability.

<https://github.com/ARM-software/mango>